

1. Assignment 4: Shadi Zidany

2. Search Trees Compared:

- AVL Tree
- Binary Search Tree (BST)
- Splay Tree

3. Key Sets Description:

Two key sets are used in the experiment:

- **Set2: $n_2 = 200,000$:** Constructed by taking a vector of integers $[0, 1, 2, \dots, (n_2 - 1)]$ and randomly shuffling it.
- **Set1: $n_1 = \text{half the size of } n_2$:** Derived from the shuffled vector v such that for each index i from 0 to $(n_1 - 1)$ the key at the i -th position is computed as $v[v[i]]$.

The keys from each set are inserted in the order they are generated.

4. Research Question:

The purpose of the experiment is to investigate how the insertion and in-order traversal performance of different search trees (self-balancing vs. non-balancing) compare when subjected to increasing dataset sizes. This comparison helps in understanding the trade-offs between faster insertions and the overhead of maintaining tree balance.

5. Program Timing Results: in milliseconds

| Operation | AVL Tree | BST | Splay Tree |
|----------------|----------|-------|------------|
| Insert set1 | 52.67 | 20.37 | 29.12 |
| Insert set2 | 97.37 | 44.45 | 65.06 |
| Traversal set1 | 1.42 | 1.74 | 2.01 |
| Traversal set2 | 3.35 | 4.02 | 4.41 |

6. Observations:

- The BST shows the fastest insertion times for both sets, while the AVL and Splay trees take longer, likely due to the additional work required for balancing.
- In-order traversal times are low across all trees, with minor variations, indicating that once built, the traversal performance is less sensitive to tree type.

7. Inferences:

The data suggest that while self-balancing trees (AVL and Splay) incur extra insertion overhead to maintain balanced structures, a standard BST—assuming the input data is random—can achieve faster insertions. However, the benefits of balanced trees may become more pronounced in scenarios involving worst-case inputs or when additional operations (like deletion or search) are considered.

8. Additional Notes:

This experiment underscores the importance of aligning data structure choice with application needs. Challenges included ensuring that the random data sets adequately tested balancing overhead. Future experiments might incorporate a broader range of key distributions (e.g., nearly sorted data) and additional operations (such as deletions) for a more comprehensive performance evaluation.